

# Socket Introduction

---

1

## Content

---

- ☐ Socket
- ☐ Stream Socket
- ☐ Datagram Socket
- ☐ APIs for managing names and IP addresses
- ☐ Socket Address Structures

---

2

## Socket

---

- ❑ What is a socket ?
- ❑ *Sockets* (in plural) are an application programming interface (API) application program and the TCP/IP stack
- ❑ A *socket* is an abstraction through which an application may send and receive data
- ❑ A socket allows an application to plug in to the network and communicate with other applications that are plugged in to the same network.

3

## Socket (cont)

---

- ❑ The main types of sockets in TCP/IP are
  - *stream sockets* : use TCP as the end-to-end protocol (with IP underneath) and thus provide a reliable byte-stream service
  - *datagram sockets* : use UDP (again, with IP underneath) and thus provide a **best-effort** datagram service
- ❑ Socket Address : include host name and port

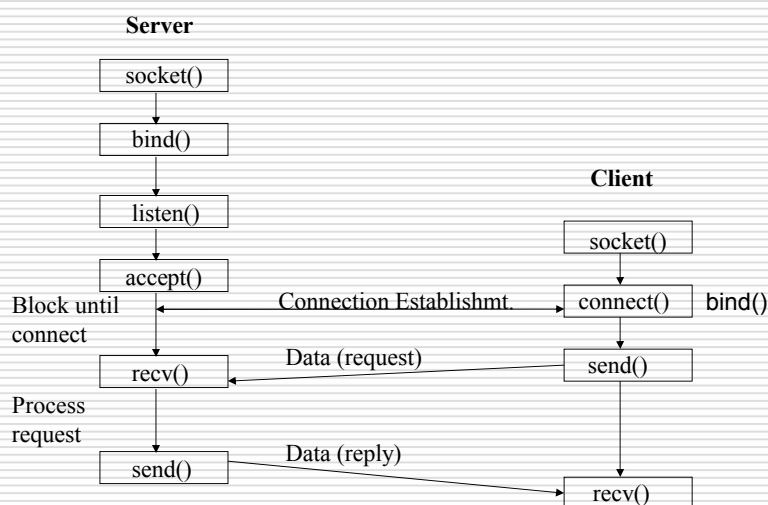
4

## Stream sockets (TCP)

- ❑ TCP provides connections between clients and servers
- ❑ TCP also provides reliability : When TCP sends data to the other end, it requires an acknowledgment in return
- ❑ TCP provides flow control
- ❑ TCP connection is full-duplex

5

## Stream sockets(TCP)



6

## Stream Socket APIs

---

- ❑ *socket*: creates a socket of a given domain, type, protocol (buy a phone)
- ❑ *bind*: assigns a name to the socket (get a telephone number)
- ❑ *listen*: specifies the number of pending connections that can be queued for a server socket. (call waiting allowance)
- ❑ *accept*: server accepts a connection request from a client (answer phone)
- ❑ *connect*: client requests a connection request to a server (call)
- ❑ *send*: write to connection (speak)
- ❑ *recv*: read from connection (listen)
- ❑ *close*: close a socket descriptor (end the call)

7

## Stream Socket APIs (cont)

---

- ❑ *socket()*
  - creates a socket of a given domain, type, protocol (buy a phone)
  - Returns a file descriptor (called a socket ID)
- ❑ *bind()*
  - Assigns a name to the socket (get a telephone number)
  - Associate a socket with an IP address and port number (Eg : 192.168.1.1:80)
- ❑ *connect()*
  - Client requests a connection request to a server
  - This is the first of the client calls

8

## Stream Socket APIs (cont)

---

### ❑ accept() :

- Server accept an incoming connection on a listening socket (request from a client)
- There are basically three styles of using accept:
  - ❑ *Iterating server*: Only one socket is opened at a time.
  - ❑ *Forking server*: After an accept, a child process is forked off to handle the connection.
  - ❑ *Concurrent single server*: use select to simultaneously wait on all open socketIds, and waking up the process only when new data arrives

9

## Stream Socket APIs (cont)

---

### ❑ listen()

- Specifies the number of pending connections that can be queued for a server socket. (call waiting allowance)

### ❑ send()

- Write to connection (speak)
- Send a message

### ❑ recv()

- read from connection (listen)
- Receive data on a socket

### ❑ close()

- close a socket (end the call)

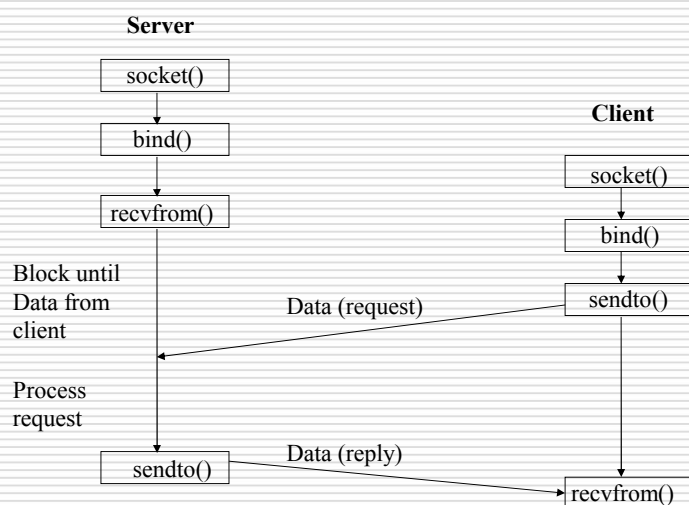
10

## Datagram Socket (UDP)

- ❑ UDP is a simple transport-layer protocol
- ❑ If a datagram is errored or lost, it won't be automatically retransmitted (can process in application)
- ❑ UDP provides a *connectionless* service, as there need not be any long-term relationship between a UDP client and server

11

## Datagram Socket (UDP)



12

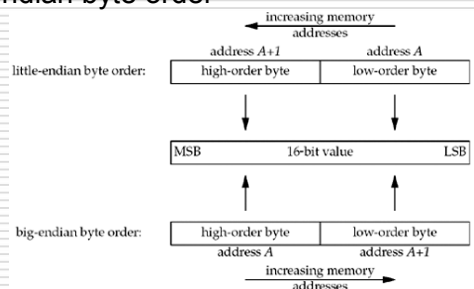
## APIs for managing names and IP addresses

- We next consider a number of auxiliary APIs:
  - *gethostname*: Returns the name of the system
  - *gethostbyname*: Get an IP address for a hostname, or vice-versa
  - *htons*, *htonl*, *ntohs*, *ntohl*: **byte ordering**
  - *inet\_ntoa*(), *inet\_aton*(), *inet\_addr*: Convert IP addresses from a dots-and-number string (eg : 192.168.1.1) to a struct *in\_addr* and back
  - *inet\_pton*, *inet\_ntop*: conversion of IP numbers between presentation and strings

13

## Byte Ordering

- There are two ways to store the two bytes in memory
  - little-endian byte order
  - big-endian byte order



14

## Byte Ordering (cont)

- ❑ There is no standard between these two byte orderings
- ❑ A variety of systems that can change between little-endian and big-endian byte ordering
- ❑ Problem : Converting between
  - *host byte order*
  - *network byte order* (The Internet protocols use big-endian byte ordering)
- ❑ Four functions to convert between these two byte orders.

15

## htons(), htonl(), ntohs(), ntohl()

- ❑ Convert multi-byte integer types from host byte order to network byte order

```
#include <netinet/in.h>
```

```
uint32_t htonl(u_long hostlong); // host to network long
uint16_t htons(u_short hostshort); // host to network short
uint32_t ntohl(u_long netlong); // network to host long
uint16_t ntohs(u_short netshort); // network to host short
```

- ❑ Each function returns the converted value.

16



## IP Number translation

---

- IP address strings to 32 bit number
- Hence, these routines translate between the address as a string and the address as the number.
- Hence, we have 4 representations:
  - IP number in host order
  - IP number in network order
  - Presentation (eg. dotted decimal)
  - Fully qualified domain name

---

17

## Socket Address Structures

---

- Most socket functions require a pointer to a socket address structure as an argument.
- Each supported protocol suite defines its own socket address structure.
- A Socket Address Structure is a structure which has information of a socket to create or connect with it
- There are three types of socket address structures
  - IPv4
  - IPv6

---

18

## IPv4 socket address structure

```
#include <netinet/in.h>
struct in_addr {
    in_addr_t s_addr; /* 32-bit IPv4 address */
                      /* network byte ordered */
};

struct sockaddr_in {
    uint8_t sin_len; /* length of structure */
    sa_family_t sin_family; /* AF_INET */
    in_port_t sin_port; /* 16-bit TCP or UDP port number */
                      /* network byte ordered */
    struct in_addr sin_addr; /* 32-bit IPv4 address */
                      /* network byte ordered */
    char sin_zero[8]; /* unused */
};
```

19

## IPv6 socket address structure

```
#include <netinet/in.h>
struct in6_addr {
    uint8_t s6_addr[16]; /* 128-bit IPv6 address */
                      /* network byte ordered */ };

#define SIN6_LEN /* required for compile-time tests */

struct sockaddr_in6 {
    uint8_t sin6_len; /* length of this struct */
    sa_family_t sin6_family; /* AF_INET6 */
    in_port_t sin6_port; /* transport layer port# */
                      /* network byte ordered */
    uint32_t sin6_flowinfo; /* flow information, undefined */
    struct in6_addr sin6_addr; /* IPv6 address */
                      /* network byte ordered */
    uint32_t sin6_scope_id; /* set of interfaces for a scope */ };
```

20

## inet\_aton()

```
#include <arpa/inet.h>
```

```
int inet_aton(const char *cp, struct in_addr *inp)
```

- ❑ Convert IP addresses from a dots-and-number string to a struct in\_addr

- ❑ Return:

- The value non-zero if the address is valid

- The value 0 if the address is invalid

```
struct in_addr someAddr;
if(inet_aton("10.0.0.1", someAddr))
    printf("The address is valid");
else printf ("The address is invalid");
```

21

## inet\_ntoa()

```
#include <arpa/inet.h>
```

```
char *inet_ntoa(struct in_addr in);
```

- ❑ Convert IP addresses from a struct in\_addr to a dots-and-number string

- ❑ Return: the dots-and-numbers string

```
struct in_addr someAddr;
if(inet_aton("10.0.0.1", someAddr))
    printf("The address is valid");
else printf ("The address is invalid");
char *addrStr;
addrStr = inet_ntoa(someAddr);
```

22

## inet\_addr()

```
#include <arpa/inet.h>
```

```
in_addr_t inet_addr(const char *cp);
```

- ❑ Convert IP addresses from a dots-and-number string to a struct `in_addr_t`

- ❑ Return:

- The value -1 if there's an error
- The address as an `in_addr_t`

```
struct in_addr someAddr;  
someAddr.s_addr = inet_addr("10.0.0.1");
```

23

## Address Resolution

24

## Content

---

- ☐ IPv4 and IPv6
- ☐ DNS
- ☐ Address and Name APIs

25

## IPv4

---

- ☐ Developed in APRANET (1960s)
- ☐ 32-bit number
- ☐ Divided into classes that describe the portion of the address assigned to the network (netID) and the portion assigned to endpoints (hosten)
  - A : netID – 8 bit
  - B : netID – 16 bit
  - C : netID – 24 bit
  - D : use for multicast
  - E : use for experiments

26

## IPv4 problem

- ❑ IPv4 addresses is being exhausted
  - ❑ Have to map multiple private addresses to a single public IP addresses (NATs)
    - Connect 2 PCs use private address space ?
    - NAT must be aware of the underlying protocols
  - ❑ IPv4 addressing is not entirely hierarchical → router must maintain routing table to deliver packets to right locations
- Develop a new version of IP Address : IPv6

27

## IPv6

- ❑ IPv6 address is 128 bits
  - To subdivide the available addresses into a hierarchy of routing domains that reflect the Internet's topology
- ❑ IPv6 address is typically expressed in 16-bit chunks displayed as hexadecimal numbers separated by colons

Example : 21DA:00D3:0000:2F3B:02AA:00FF:FE28:9C5A

or : 21DA:D3:0:2F3B:2AA:FF:FE28:9C5A

28

## DNS (Domain Name System)

---

- ❑ Computers use IP Addresses to connect hosts
  - What about humans ? – IP Addresses are very complex and hard to remember (for people)
- ❑ Use name instead of IP Address → Domain Name System
- ❑ Problem of DNS
  - People use names, Computers use IP Addresses → translate between two spaces
  - Domain name system must be hierarchical (for management and maintain)
- ❑ Domain name space : divide to zones

29

## DNS (cont)

---


- ❑ How to translate between domain name-IP Address and reverse ?
  - DNS Resolver
  - DNS Server
- ❑ A DNS query
  - A *non-recursive query* : DNS server provides a record for a domain for which it is authoritative itself, or it provides a partial result without querying other servers
  - A *recursive query* : DNS server will fully answer the query by querying other name servers
- ❑ DNS primarily uses User Datagram Protocol (UDP) on port number 53 to serve requests

30

## DNS (cont)

### □ Address resolution mechanism

- Local system is pre-configured with the known addresses of the root server in a file of *root hints*
- Query one of the root servers to find the server authoritative for the next level down
- Querying level down server for the address of a DNS server with detailed knowledge of the lower level domain until reach the DNS Server return final address



31

## DNS (cont)

- A *Resource Record* (RR) is the basic data element in the domain name system
- All records use the common format specified in RFC 1035 (in IP networks)
- **RR (Resource record) fields**
  - NAME (variable)
    - Name of the node to which this record pertains.
  - TYPE (2)
    - Type of RR. For example, MX is type 15
  - CLASS (2)
    - Class code
  - TTL (4)
    - Unsigned time in seconds that RR stays valid
  - RDLENGTH (2)
    - Length of RDATA field
  - RDATA (variable)
    - Additional RR-specific data

32



## List of Address and Name APIs

---

#include <sys/socket.h>

- ❑ **gethostbyaddr()**
  - Retrieve the name(s) and address corresponding to a network address.
- ❑ **gethostname()**
  - Retrieve the name of the local host.
- ❑ **gethostbyname()**
  - Retrieve the name(s) and address corresponding to a host name.
- ❑ **getprotobyname()**
  - Retrieve the protocol name and number corresponding to a protocol name.
- ❑ **getprotobynumber()**
  - Retrieve the protocol name and number corresponding to a protocol number.
- ❑ **getservbyname()**
  - Retrieve the service name and port corresponding to a service name.
- ❑ **getservbyport()**
  - Retrieve the service name and port corresponding to a port.

---

33

## New APIs for IPv6

---

- ❑ Those APIs only supports IPv4 but IPv6 will be replace IPv4 in the future, so we need APIs support IPv6
- ❑ They are
  - getaddrinfo
  - getnameinfo
- ❑ These APIs have replaced the IPv4 specific routines

---

34

## gethostbyaddr()

- Get host information corresponding to an address.

```
#include <netdb.h>
#include <sys/socket.h>
struct hostent *gethostbyaddr (in_addr *addr, socklen_t len, int family);
```

- addr
  - A pointer to an address in network byte order.
- len
  - The length of the address, which must be 4 for PF\_INET addresses.
- type
  - The type of the address, which must be PF\_INET.
- Return value
  - If no error occurs, gethostbyaddr() returns a pointer to the hostent structure
  - Otherwise it returns a NULL pointer and a specific error number

35

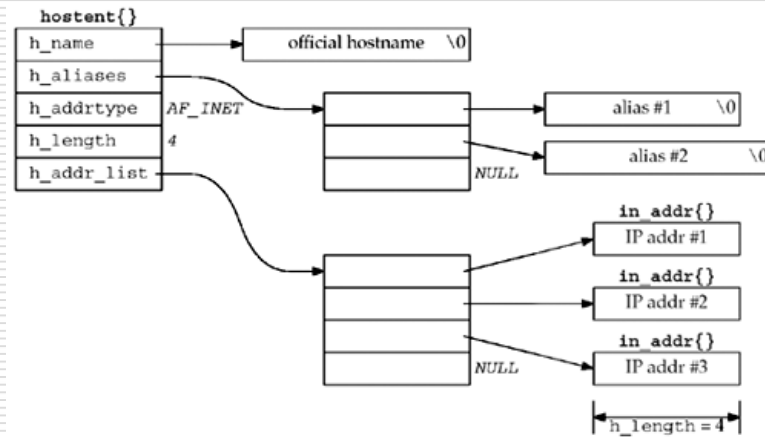
## struct hostent

```
struct hostent {
    char *h_name;      /* official (canonical) name of host */
    char **h_aliases;  /* pointer to array of pointers to
                        alias names */
    int h_addrtype;    /* host address type: AF_INET */
    int h_length;      /* length of address: 4 */
    char **h_addr_list; /* ptr to array of ptrs with IPv4 addrs */
};
```

- what is this struct hostent that gets returned?
- It has a number of fields that contain information about the host in question.

36

## struct hostent



37

## gethostname()

- ❑ Return the standard host name for the local machine.

```
#include <sys/unistd.h>
#include <sys/socket.h>
int gethostname(char *name, size_t len);
```

- ❑ name
  - A pointer to a buffer that will receive the host name.
- ❑ Len
  - The length of the buffer
- ❑ Return value
  - If no error occurs, `gethostname()` return 0
  - Otherwise it returns `SOCKET_ERROR` and a specific error code

38

## gethostbyname()

- Get host information corresponding to a hostname.

```
#include <netdb.h>
#include <sys/socket.h>
struct hostent *gethostbyname (const char *hostname);
```

- name
  - A pointer to the name of the host
- Returns a pointer to a hostent structure as described under gethostbyaddr().
- Return value
  - If no error occurs, gethostbyname() returns a pointer to the hostent structure described above.
  - Otherwise it returns a NULL pointer and a specific error number

39

## getservbyname()

- Get service information corresponding to a service name and protocol.

```
#include <netdb.h>
#include <sys/socket.h>
struct servent *getservbyname (const char *servname, const
                                char *proto name);
```

- servname
  - A pointer to a service name.
- proto name
  - An optional pointer to a protocol name.
  - If this is NULL, getservbyname() returns the first service entry for which the name matches the s\_name or one of the s\_aliases.
  - Otherwise getservbyname() matches both the name and the proto.
- Returns
  - non-null pointer if OK
  - NULL on error

```
struct servent *sptr;
sptr = getservbyname("ftp", "tcp");
```

40

## struct servent

```
struct servent {
    char *s_name;        /* official service name */
    char **s_aliases;    /* alias list */
    int s_port;          /* port number, network-byte order */
    char *s_proto;       /* protocol to use */
};
```

- s\_name
  - Official name of the service.
- s\_aliases
  - A NULL-terminated array of alternate names.
- s\_port
  - The port number at which the service may be contacted. Port numbers are returned in network byte order.
- s\_proto
  - The name of the protocol to use when contacting the service.

41

## getservbyport

- Get service information corresponding to a port and protocol.

```
#include <netdb.h>
#include <sys/socket.h>
struct servent *getservbyport (int port, const char *proto name);
```

- port
  - The port for a service, in network byte order.
- proto name
  - An optional pointer to a protocol name.
  - If this is NULL, getservbyport() returns the first service entry for which the port matches the s\_port.
  - Otherwise getservbyport() matches both the port and the proto.
- Return
  - non-null pointer if OK
  - NULL on error

```
struct servent *sptr;
sptr = getservbyport (htons (53), "udp");
```

42